

The Platform Design Problem

Christos Papadimitriou, **Kiran Vodrahalli**, Mihalis Yannakakis

[Columbia University](#)

Google Learning Theory Reading Group

October 11, 2021

The Data-Collection Problem


- Modern machine learning requires large amounts of high-quality data
- Collecting supervised labels is expensive
- Unsupervised learning is challenging to use
- Is it possible to create environments which generate useful data?
 - Ex: Reddit users provide sarcasm labels using the “/s” tag

The Data-Collection Problem

- Modern machine learning requires large amounts of high-quality data
 - Collecting supervised labels is expensive
 - Unsupervised learning is challenging to use
- Is it possible to create environments which generate useful data?
 - Ex: Reddit users provide sarcasm labels using the “/s” tag

Modern tech companies try to solve this problem.

Economics of the Online Firm



- User data feeds revenue
 - Better demand segmentation
 - Ad/recommendation revenue
 - Better models => better services
- Online services bring value
 - Convenience
 - Knowledge

Outline

- Problem Definition
- General Case
- Tractable “Flower” Case
 - Agent Behavior
 - Designer’s Algorithm
- Extensions
- Summary
- Future Work

Problem Definition

Platform Design

Problem

Model the revenue-maximization problem of today's online firms (e.g. Google, FB, etc.) and understand computational tractability.

Bi-Level MDP Optimization Model

Agent: participates in Life MDP

Designer: tweaks the Life MDP by building platforms.

Goal: **Designer** wants to indirectly optimize its reward via **Agent**'s optimal behavior! (Find Stackelberg)

- Key Idea: Google builds various apps (Maps, Search, Social Network, etc.) and profits based on usage of these apps.
- The usage of apps modifies the transitions of the Markov Chain of the user's life
- Assume the Designer has linear rewards over the steady state distribution of the resulting Markov chain (agent policy + Life MDP)

The Stackelberg Game

- Designer moves first:
 - Adds **platforms** which, if adopted, modify transitions to an existing Markov Chain
- Agent moves second:
 - Receives **MDP** from Designer, plays optimal behavior
- Example of bi-level MDP optimization
- What is the computational complexity of solving for equilibrium?

Formal Problem Statement

- An **agent** lives in an irreducible Markov chain with $A = [n]$ states.
- The **designer** chooses $S \subseteq A$ states to add platforms to.
- The agent may **adopt or not adopt** the platform at each state:
 - If **adopt**, the transitions change. Otherwise they do not.
 - Assume the chain remains irreducible.


Formal Problem Statement

- Assign a utility rate for the agent (c_i) and the designer (d_i) at $i \in [n]$.
- The agent solves the resulting Markov Decision Process.
 - Resulting steady-state probabilities are given by π .
- The designer optimizes over S :

$$\text{profit}(S) := \sum_{i \in S} d_i \cdot \pi_i(S) - \sum_{i \in S} \text{cost}_i$$

General Case


Picture of the General Case



What platforms
should I build?



Picture of the General Case




What platforms
should I build?



At a cost, the firm can **add an opt-in action** to platforms they create (ex: Google Maps).

Picture of the General Case



Maybe we
should create
Maps
technology....



Online firm

Builds platform
Maps at a cost.


Opt in to Maps

Computational Tractability I: General Case


- It is **strongly NP-hard** to decide whether the Designer can obtain positive profit – and therefore **hard to approximate**.
- Reduction from SET COVER
 - Designer builds platforms which each solve subset of Agent's problems.
 - Most cost-effective covering set is NP hard.
- In economic terms, the reduction exploits the complexity of "**complementary goods**."
 - Ex: Brick-and-mortar retail ads help the Agent discover the store, Maps helps the Agent get to the store.

Tractable “Flower” Case

A More Tractable Case: The Flower



Life MDP



Tweaked MDP via y_i

A More Tractable Case: The Flower

- Problem can be solved by an FPTAS
- Why tractable?
 - Substitutes rather than complements
 - Allocate time spent in each platform
 - Simpler low-level behavior (greedy agent)
 - Admits a DP upon discretization (knapsack DP)

Agent Behavior

The Agent's Greedy Algorithm

- Solving for the steady state distribution yields a quasi-concave combinatorial optimization problem:

Lemma 1. *The agent's objective for an optimal policy defined in Section 2 can be re-written as the following optimization in the special case of the flower MDP (Definition 2):*

$$\operatorname{argmax}_{S \subseteq [n]} \frac{A + \sum_{j \in S} z_j \phi(j)}{B + \sum_{j \in S} z_j} \quad (1)$$

where

$$A := \sum_{i=1}^n \lambda_i c_i^{\text{life}}; \quad B := 1 + \sum_{i=1}^n \lambda_i; \quad \lambda_i = \frac{p_i}{1 - q_i}; \quad z_i = \frac{p_i}{1 - q_i - y_i} - \frac{p_i}{1 - q_i} \geq 0;$$

$$\phi(i) := \begin{cases} c_i^{\text{platform}} + \frac{\lambda_i}{z_i} (c_i^{\text{platform}} - c_i^{\text{life}}) & \text{if } z_i > 0 \\ 0 & \text{if } z_i = 0 \end{cases};$$

We therefore define

$$\text{utility}^{\text{Agent}}(S) := \frac{A + \sum_{j \in S} z_j \phi(j)}{B + \sum_{j \in S} z_j}$$

The Agent's Greedy Algorithm

- Solving for the steady state distribution yields quasi-concave combinatorial optimization problem:

Lemma 1. *The agent's objective for an optimal policy defined in Section 2 can be re-written as the following optimization in the special case of the flower MDP (Definition 2):*

$$\operatorname{argmax}_{S \subseteq [n]} \frac{A + \sum_{j \in S} z_j \phi(j)}{B + \sum_{j \in S} z_j} \quad (1)$$

where

$$A := \sum_{i=1}^n \lambda_i c_i^{\text{life}}; \quad B := 1 + \sum_{i=1}^n \lambda_i; \quad \lambda_i = \frac{p_i}{1 - q_i}; \quad z_i = \frac{p_i}{1 - q_i - y_i} - \frac{p_i}{1 - q_i} \geq 0;$$

$$\phi(i) := \begin{cases} c_i^{\text{platform}} + \frac{\lambda_i}{z_i} (c_i^{\text{platform}} - c_i^{\text{life}}) & \text{if } z_i > 0 \\ 0 & \text{if } z_i = 0 \end{cases}$$

Potential function

We therefore define

$$\text{utility}^{\text{Agent}}(S) := \frac{A + \sum_{j \in S} z_j \phi(j)}{B + \sum_{j \in S} z_j}$$

The Agent's Greedy Algorithm

ALGORITHM 1: GREEDY ALGORITHM

Input: Parameters of the Agent's problem: transition probabilities and utility coefficients in and out of the platform.

Output: An optimal subset $S \subseteq [n]$ of states where the Agent accepts the platform.

Initialize $S := \{\}$

for $k \in [n]$ sorted⁹ from largest to smallest $\phi(k)$ **do**

if $\text{utility}^{\text{Agent}}(S) < \phi(k)$ **then**

 Update $S := S \cup \{k\}$

else

return S

end

end

return S

Sort states by potential function and add until utility = potential

Proof Outline

- Application of the mediant inequality: $\frac{x}{y} < \frac{r}{s} \iff \frac{x}{y} < \frac{x+r}{y+s} < \frac{r}{s}$ where $y, s > 0$.
- Optimal policy must be “prefix”: must contain first m states, sorting by potential function ϕ .
- Proof and algorithm slightly more complex when $z_i < 0$.
- ϕ is the platform reward + scaled gain over regular life.
- Agent accepts when
 - Base platform reward is high
 - Amount of time spent on platform is high and platform is better

Designer's Algorithm

Recall: Designer's Objective

$$\text{profit}(S) := \sum_{i \in S} d_i \cdot \pi_i(S) - \sum_{i \in S} \text{cost}_i$$

Set of states to build platforms

Designer's steady-state reward rates

Agent's steady state probabilities


Designer's one-time costs for building each platform

Restrictions

- Expanding the profit function given agent behavior:

$$\text{profit}(S) := \frac{\sum_{i \in \text{Agent}(S)} d_i \cdot \frac{p_i}{1-q_i-y_i}}{B + \sum_{i \in \text{Agent}(S)} z_i} - \sum_{i \in S} \text{cost}_i$$

$P_1(S)$ $D(S)$



- Define $\max_i d_i =: K$
 - Maximum profit is nK
- Assume z_i are $\text{poly}(n)$ and discretized with gap δ and costs are $K * \text{poly}(n)$

Target Algorithm

- Deciding whether it is possible to attain a certain profit is NP complete
- Reduction from PARTITION
- Thus, our goal: A $(1 - \epsilon)$ approximate algorithm in polynomial time.

The Designer's Dynamic Program

- **Key Idea:** Use a (poly-sized) hash table with rounded rewards
- Difficulty comes from profit scale and non-discretized z_i
- Hash function:

$$\text{hash}(S) := \left(\lceil \frac{\text{profit}(S)}{\epsilon K / 2n} \rceil, \lceil \frac{P_1(S)}{\epsilon K / 2n} \rceil, \mathbf{D}(S) / \delta \right)$$

- Similar to standard Knapsack FPTAS (Ibarra & Kim, 1975)

The Designer's Dynamic Program

ALGORITHM 2: DESIGNER'S FPTAS FOR THE PDP

Input: The parameters of the PDP: transition probabilities, utility and cost coefficients for the Agent and the Designer, and small positive reals ϵ, δ

Output: A $(1 - \epsilon)$ -approximately optimal subset of states S^* for which to deploy platforms.

$\mathbf{N}(S)$ and $\mathbf{D}(S)$ denote the numerator and the denominator of the Agent's objective function, with the constant terms omitted

$P_1(S)$ denotes the first term in the Designer's profit function

SET is a hash table of subsets of $[n]$ indexed by triples of integers

The hash function is $\text{hash}(S) := \left(\lceil \frac{\text{profit}(S)}{\epsilon K/2n} \rceil, \lceil \frac{P_1(S)}{\epsilon K/2n} \rceil, \mathbf{D}(S)/\delta \right)$

Initialize the hash table SET to contain only the empty set in the bin $(0, 0, 0)$

for $k \in [n]$ **do**

for $S \in \text{SET}$ in lexicographic order **do**

$S' := S \cup \{k\}$

if *Agent will adopt all platforms in S' and $\text{profit}(S') > 0$* **then**

if $\text{hash}(S') \in \text{SET}$ **then**

$\hat{S} := \text{SET}[\text{hash}(S')]$

if $\mathbf{N}(\hat{S}) > \mathbf{N}(S')$ **then**

$\text{SET}[\text{hash}(S')] := S'$

end

else

$\text{SET}[\text{hash}(S')] := S'$

end

end

end

return *the set S in the hash table with largest first hash value*

efficient feasibility check

update if numerator of agent reward is smaller
than current hash

new hash

Proof Outline: Key Lemma

- For S, S' that hash to same bin; if $N(S) \leq N(S')$, for any postfix set T :
 - If $S' \cup T$ is feasible, so is $S \cup T$
 - $S \cup T$ is at most $\epsilon K / n$ worse than $S' \cup T$

Proof:

- Feasibility: Shared denominator + sub-optimality of $S \cup T$
- Suboptimality: $|\text{profit}(S \cup T) - \text{profit}(S' \cup T)| \leq |\text{profit}(S) - \text{profit}(S')| + |P_1(S) - P_1(S')|$ and shared hash bin.

Proof Outline: Applying Key Lemma

- After i^{th} iteration:
 - hash table contains S extendable to $S \cup T$ at most ϵ^{iK}/n suboptimal.
- After n iterations, suboptimal by at most $(1 - \epsilon)$ factor.

Proof: By induction.

- Inductive step: apply the Key Lemma at step i to get:
 - If S is replaced this step: $\epsilon^{(i-1)K}/n + \epsilon^K/n = \epsilon^{iK}/n$ suboptimal
 - Otherwise: even tighter bound, no degradation.

Proof Outline: Complexity

- Checking feasibility is $O(n \log n)$
- Three dimensions of hash table:
 - Denominator dimension is size n/δ
 - Profit dimensions are each size $\frac{nK}{\epsilon K/n} = \frac{n^2}{\epsilon}$ times a possible polynomial factor for costs
- Total size of hash table: $O(\text{poly}(n) \cdot n^5 / \epsilon^2 \delta)$
- Thus: polynomial runtime.

Extensions

Multiple Agents

- Replace designer objective with summation over agents:


$$\text{profit}(S) := \sum_i \frac{\sum_{j \in \text{Agent}_i(S)} d_{ij} \cdot \frac{p_{ij}}{1 - q_{ij} - y_{ij}}}{B_i + \sum_{l \in \text{Agent}_i(S)} z_{il}} - \sum_{j \in S} \text{cost}_j$$

- An exact polytime DP exists if #agents is constant.
 - Exponential in #agents
 - Also require potentials ϕ_i to be discretized by δ' with poly size.
- No FPTAS for 2 agents if ϕ_i not polynomial size.

Multiple Agents

- **Key ideas:**
 - Discretize over potentials and denominators.
 - For each potential-denominator pair (θ, D) , compute optimal subset s.t. potentials are at least θ with a DP hash table (size M^{3k}).
 - Enumerate over (θ, D) to get global optimum.
- **Proof sketch:** Key idea is that we can exactly compute values for each entry in the (θ, D) hash table.

Designer Competition



What platforms
should I build to
compete?



Multiple Platforms (Flower Setting)

- What if other competing designers have already built platforms?
 - Each platform affects only one state
 - At most one for each designer per state
- How does an agent behave?
- How should a designer optimally place platforms?


Multiple Platforms (Flower Setting)

- Agent's algorithm is still greedy – but different potential function
- For platforms j, j' at the same state, define:

$$\rho(j, j') = \frac{z_{j'}\phi(j') - z_j\phi(j)}{z_{j'} - z_j}$$

Multiple Platforms (Flower Setting)

- Each state's non-redundant platforms obey the following structure:



Multiple Platforms (Flower Setting)

- **Proof idea:** Greedy swap argument with new potential, many details.
- The new potential function:

$$\psi(j, \ell) := \begin{cases} \phi(j, 1) & \text{if } \ell = 1 \\ \rho_j(\ell, \ell - 1) & \text{if } \ell > 1 \end{cases}$$

state Platform
index, sorted
by ϕ in
decreasing
order

Multiple Platforms (Flower Setting)

ALGORITHM 4: MULTI-PLATFORM AGENT'S ALGORITHM

Input: Parameters of the Agent's problem: transition probabilities and utility coefficients in and out for all platforms.

Output: An optimal feasible subset S of platforms.

Remove redundant platforms for each state

Compute the parameters ψ for the (nonredundant) platforms

Sort the platforms in decreasing order $\psi(j)$

Initialize $S := \{\}$

for each platform j in decreasing order of $\psi(j)$ **do**

if $\psi(j) \leq u(S)$ **then**

return S

 Agent utility

else

if j is the first platform for its state **then**

 Update $S := S \cup \{j\}$

else

 Update $S := S \cup \{j\} \setminus \{prev(j)\}$

end

end

 Replace the platform at state j

end

return S

Runs in time
 $O(n + m \log m)$

$n = \#$ states

$m = \#$ total platforms

Multiple Platforms (Flower Setting)

- Is there an efficient designer algorithm?
- The multi-agent algorithm also (essentially) works in the multi-platform setting
 - Same discretization assumptions (potentials, denominator)
 - Exact algorithm
 - Polynomial time when #agents is constant
- Slight difference from old algorithm:
 - Modify the hash function: numerator and denominator of ψ instead

Summary

Recap

- Platform design: model economic activity of online firms
- General case of platform design is strongly NP complete.
- Tractable special case: the flower MDP
- Greedy agent algorithm
- Knapsack-style DP FPTAS for designer w/unbounded potentials
- Under polynomial, discretized potentials, exact DP for k agents ($\text{poly}(n) \cdot 2^k$)
- Similar for multiple platforms
- **Many open directions!**

Future Work

Future Work

- Designer vs. designer
 - Complexity of pure Nash
 - Repeated game settings
- Privacy/fairness questions for agent
- Other classes of tractable MDPs?
- Results for generic classes of agent behavior?
- Many questions are problems of formulation

Learning Theoretic Questions

- What if agent/designer have to learn?

Learning Theoretic Questions

- What if agent/designer have to learn?
- Optimizing over distribution of agent types w/finite support
 - Expected reward: smoothed version of the objective
- ERM with quasi-linear function class
 - We can solve ERM efficiently if finite support is constant (with discretized, poly-bounded potentials)

Learning Theoretic Questions

- What if agent/designer have to learn?
- Optimizing over distribution of agent types w/finite support
 - **Open:** What if support isn't constant size, or is continuous?
 - ERM via our algorithm no longer computationally efficient
 - Other approaches? Under what conditions is computationally efficient learning possible?
 - Goal: beat $O(2^k \text{poly}(n))$ algorithms for k -type supports

Learning Theoretic Questions

- What if agent/designer have to learn?
- Combinatorial bandit setting
 - Suppose the designer is an online bandit
 - Plays combinatorial set S each round
 - **Open: Complicated dependencies between arms + nonlinear rewards**

Learning Theoretic Questions

- What if agent/designer have to learn?
- Repeated game variants of the problem where both agent and designer are learners
- Other equilibria: what is computationally tractable?
- Strategic agents and designer-designer competition

Learning Theoretic Questions

- What if agent/designer have to learn?
- Remove the abstraction of designer rewards:
 - Agents emit data distributions
 - Restricted sampling conditions
 - Goal: Solve some learning problem about agents
 - Connected with data valuation
 - How to design sampling environment?

Grand Vision

- **Design environments** which generate useful, sampleable data
- **Model economics** of companies dependent on information economy
- **Model strategic behavior** of online firms and their users
- **Reinforcement learning** aided by environment design
- **Manipulation and resistance** of learning agents